

Assignment 4: Logistic Regression and Gradient Descent

Machine Learning

Fall 2021

🔗 Learning Objectives

- Learn about the logistic regression algorithm.
- Learn about gradient descent for optimization.

🔗 Prior Knowledge Utilized

- Supervised learning problem framing.
- Log loss

Imagine that instead of our model, \hat{f} , spitting out either 0 or 1, it outputs a confidence that the input \mathbf{x} has an output $y = 1$. In other words, rather than giving us its best guess (0 or 1), the classifier would indicate to us its degree of certainty regarding its prediction. This notion of “certainty” can be formalized using the concept of a probability.

If our model outputs a probability p when supplied with an input \mathbf{x} (i.e., $\hat{f}(\mathbf{x}) = p$), we might then ask ourselves what loss function we should choose in order to select the best possible model? This loss function will be used to quantify how bad a prediction p is given the actual output y (recall that for binary classification the output is either 0 or 1).

1. **0-1 loss:** we will extract from the weatherperson’s prediction the most likely output (e.g., if $p = 0.75$, that would be sunny, if $p = 0.4$, that would be rainy). If the most likely output matches the actual output we give a loss of 0, otherwise we give a loss of 1 (this is similar to Equation ??).
2. **squared loss:** one downside of 0-1 loss is that it doesn’t take into account the certainty expressed by the weatherperson. The weatherperson gets the same loss if it is rainy and they predicted $p = 0.51$ or $p = 1$. For squared loss we compute the difference between the outcome and p and square it to arrive at the loss. For example if the weatherperson predicts $p = 0.51$ and it is sunny the loss is $(1 - 0.51)^2$. If it was rainy in this same example, the loss is $(0 - 0.51)^2$.

1 The Classification Problem

So far in this class we’ve looked at supervised learning problems where the responses y_i are continuous-valued and the loss function is quadratic ($\ell(y, \hat{y}) = (y - \hat{y})^2$). This is an example of a regression problem. There are many times, however, where it is unnatural to frame a problem as a regression. For instance, it may be the case that y_i does not come from a continuous range but instead can only take on a few different values. This sort of problem is known as a classification problem. For instance, you might want to

have a system that takes in an image of a person and predicts their identity. The identity could be thought of as the output, y_i , and it would only make sense for y_i to be one of several values (e.g., each value might represent a particular person the system was trained to recognize). In this assignment you'll learn about a special case of the classification problem known as binary classification (where y_i is either 0 or 1, e.g., a Paul versus Sam recognizer).

In this assignment we will formalize the binary classification problem and see a very useful algorithm for solving it called *logistic regression*. You will also see that the logistic regression algorithm is a very natural extension of linear regression. Our plan for getting there is going to be pretty similar to what we did for linear regression.

- Build some mathematical foundations
- Introduce logistic regression from a top-down perspective
- Learn about logistic regression from a bottom-up perspective

2 Logistic Regression (top-down)

Now that we have built up some understanding of how probabilities can be used as a way of quantifying confidence in predictions, you are ready to learn about the logistic regression algorithm.

As always, we assume we are given a training set of inputs and outputs. As in linear regression we will assume that each of our inputs is a d -dimensional vector \mathbf{x}_i and since we are dealing with binary classification, the outputs, y_i , will be binary numbers (indicating whether the input belongs to class 0 or 1). Our hypothesis functions, \hat{f} , output the probability that a given input has an output of 1. What's cool is that we can borrow a lot of what we did in the last couple of assignments when we learned about linear regression. In fact, all we're going to do in order to make sure that the output of \hat{f} is between 0 and 1 is pass $\mathbf{w}^\top \mathbf{x}$ through a function that “squashes” its input so that it outputs a value between 0 and 1. This idea is shown graphically in Figure 1.

To make this intuition concrete, we define each \hat{f} as having the following form (note: this equation looks daunting. We have some tips for interpreting it below).

$$\begin{aligned} \hat{f}(\mathbf{x}) &= \text{probability that output, } y, \text{ is } 1 \\ &= \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}}} \end{aligned} \quad (1)$$

Here are a few things to notice about this equation:

1. The weight vector that we saw in linear regression, \mathbf{w} , has made a comeback. We are using the dot product between \mathbf{x} and \mathbf{w} (which creates a weighted sum of the x_i 's), just as we did in linear regression!
2. As indicated in Figure 1, the dot product $\mathbf{w}^\top \mathbf{x}$ has been passed through a squashing function known as the [sigmoid function](#). The graph of $\sigma(u) = \frac{1}{1+e^{-u}}$ is shown in Figure 2. $\sigma(\mathbf{w}^\top \mathbf{x})$ is exactly what we have in Equation 1.




Figure 1: Graphical representation of both linear and logistic regression. The key difference is the application of the squashing function shown in yellow. [Original source](#).




Figure 2: a graph of the sigmoid function $\frac{1}{1+e^{-x}}$.

⚠ Notice

Pretty much every writeup of logistic regression contains a simple two-independent variable example. Usually these examples involve things like predicting who would be approved for a credit card or who will be admitted to a college. We are going to do something a little messier and bit different. If you want the admission to college example, you can read about it in [Building a Logistic Regression in Python](#).

3 Putting it All Together: Deriving the Logistic Regression Learning Rule

Let's summarize what we've done thus far in this assignment.

- We motivated the binary classification problem.
- We presented a particular useful loss function (log loss).
- We met the logistic regression model and tried it out on a real dataset.

Next, we're going to build on these pieces and formalize the logistic regression problem and derive a learning rule to solve it (i.e., compute the optimal weights). The formalization of logistic regression will combine Equation 1 with the selection of ℓ to be log loss. This choice of ℓ results in the following objective function.

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_{i=1}^n \left(-y_i \ln \sigma(\mathbf{w}^\top \mathbf{x}_i) - (1 - y_i) \ln(1 - \sigma(\mathbf{w}^\top \mathbf{x}_i)) \right) \quad (2)$$

$$= \arg \min_{\mathbf{w}} \sum_{i=1}^n \left(-y_i \ln \left(\frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) - (1 - y_i) \ln \left(1 - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \right) \quad \text{expanded out if you prefer this form} \quad (3)$$

While this looks a bit crazy, since y_i is either 0 or 1, the multiplication of the expressions in the summation by either y_i or $1 - y_i$ are essentially acting like a switch—depending on the value of y_i we either get one term or the other. Our typical recipe for finding \mathbf{w}^* has been to take the gradient of the expression inside the arg min, set it to 0, and solve for \mathbf{w}^* (which will be a critical point and hopefully a minimum). The last two steps will be a bit different for reasons that will become clear soon, but we will need to find the gradient. We will focus on finding the gradient in the next couple of parts.

3.1 Useful Properties of the Sigmoid Function

Looking at Equation 3 it looks really, really hairy! We see that in order to compute the gradient we will have to compute the gradient of $\mathbf{x}^\top \mathbf{w}$ with respect to \mathbf{w} (we just wrapped our minds around this last assignment). Additionally, we will have to take into account how the application of the sigmoid function and the log function changes this gradient. In this section we'll learn some properties for manipulating the sigmoid function and computing its derivative.

Exercise 1 (60 minutes)

The sigmoid function, σ , is defined as

$$\sigma(x) = \frac{1}{1 + e^{-x}} . \quad (4)$$

- (a) Show that $\sigma(-x) = 1 - \sigma(x)$.
- (b) Show that the derivative of the logistic function $\frac{d}{dx}\sigma(x) = \sigma(x)(1 - \sigma(x))$

3.2 Chain Rule for Gradients

We now know how to take derivatives of each of the major pieces of Equation 3. What we need is a way to put these derivatives together. You probably remember that in the case of single variable calculus you have just such a tool. This tool is known as the chain rule. The chain rule tells us how to compute the derivative of the composition of two single variable functions f and g .

$$\begin{aligned} h(x) &= g(f(x)) & h(x) \text{ is the composition of } f \text{ with } g \\ h'(x) &= g'(f(x))f'(x) & \text{this is the chain rule!} \end{aligned} \quad (5)$$

Suppose that instead of the input being a scalar x , the input is now a vector, \mathbf{w} . In this case h takes a vector input and returns a scalar, f takes a vector input and returns a scalar, and g takes a scalar input and returns a scalar.

$$\begin{aligned} h(\mathbf{w}) &= g(f(\mathbf{w})) & h(\mathbf{w}) \text{ is the composition of } f \text{ with } g \\ \nabla h(\mathbf{w}) &= g'(f(\mathbf{w}))\nabla f(\mathbf{w}) & \text{this is the multivariable chain rule} \end{aligned} \quad (6)$$

Exercise 2 (60 minutes)

- (a) Suppose $h(x) = \sin(x^2)$, compute $h'(x)$ (x is a scalar so you can apply the single-variable chain rule).
- (b) Define $h(\mathbf{v}) = (\mathbf{c}^\top \mathbf{v})^2$. Compute $\nabla_{\mathbf{v}} h(\mathbf{v})$ (the gradient of the function with respect to \mathbf{v}).
- (c) Compute the gradient of the expression from Equation 3 (reproduced below for your convenience).

$$\sum_{i=1}^n -y_i \ln \sigma(\mathbf{w}^\top \mathbf{x}_i) - (1 - y_i) \ln (1 - \sigma(\mathbf{w}^\top \mathbf{x}_i)) . \quad (7)$$

You can either use the chain rule and the identities you learned about sigmoid, or expand everything out and work from that.

3.3 Gradient Descent for Optimization

If we were to follow our derivation of linear regression we would set our expression for the gradient to 0 and solve for \mathbf{w} . It turns out this equation will be difficult to solve due to the σ function. Instead, we can use an iterative approach where we start with some initial value for \mathbf{w} (we'll call the initial value \mathbf{w}^0 , where the superscript corresponds to the iteration number) and iteratively adjust it by moving down the gradient (the gradient represents the direction of fastest increase for our function, therefore, moving along the negative gradient is the direction where the loss is decreasing the fastest).

External Resource(s) (45 minutes)

There are tons of great resources that explain gradient descent with both math and compelling visuals.

- Recommended: [Gradient descent, how neural networks learn | Deep learning, chapter 2](#) (start at 5:20)
- An Introduction to Gradient Descent ([on NB](#), [original](#))
- The Wikipedia page on Gradient Descent ([on NB](#), [original](#))
- [Ahmet Sacan's video on gradient descent](#) (this one has some extra stuff, but it's pretty clearly explained).
- There are quite a few resources out there, do you have some suggestions? (suggest so on NB)

Exercise 3 (10 minutes)

To test your understanding of these resources, here are a few diagnostic questions.

- When minimizing a function with gradient descent, which direction should you step along in order to arrive at the next value for your parameters?
- What is the learning rate and what role does it serve in gradient descent?
- How do you know when an optimization performed using gradient descent has converged?
- True or false: provided you tune the learning rate properly, gradient descent guarantees that you will find the global minimum of a function.

If we take the logic of gradient descent and apply it to the logistic regression problem,

we arrive at the following learning rule. Given some initial weights \mathbf{w}^0 , and a learning rate η , we can iteratively update our weights using the formula below.

$$\mathbf{w}^{n+1} = \mathbf{w}^n - \eta \sum_{i=1}^n -(y_i - \sigma(\mathbf{w}^\top \mathbf{x}_i)) \mathbf{x}_i \quad \text{applying the result from exercise 4} \quad (8)$$

$$= \mathbf{w}^n + \eta \sum_{i=1}^n (y_i - \sigma(\mathbf{w}^\top \mathbf{x}_i)) \mathbf{x}_i \quad \text{distribute the negative} \quad (9)$$

This beautiful equation turns out to be the recipe we need to implement logistic regression.

Exercise 4 (10 minutes)

Now that you've had a second round of working with an in-class partner, we want to hear more about how it's going.

- (a) Did you and your partner work well together? To help us understand how this better, provide a concrete examples to illustrate what went well and what didn't. If you want to avoid working with you partner in the future, say so (make sure to include the person's name). Only use this option in extenuating circumstances.
- (b) We are assigning partners as an experiment. Our rationale for this choice is that finding a partner is potentially stressful and it may be difficult to work with folks you don't normally socialize with. With this rationale in mind, do you think we should continue to assign class work partners? Should we tweak this somehow?